

Contents

Confluent Platform GitOps with ArgoCD	1
Complete Guide to Declarative Kafka Management on Kubernetes	1
Table of Contents	1
Executive Summary	2
Architecture Overview	2
Prerequisites	4
ArgoCD Installation & Configuration	5
Helm Charts Structure	6
Deploying Confluent Platform	7
Deploying Kafka Streams Applications	8
Deploying Flink Applications	9
Configuration Management	11
Branch-Based Deployments	11
Scaling Applications	12
Secrets Management	13
CI/CD Integration	14
Monitoring & Observability	15
Operational Commands Reference	16
Troubleshooting Guide	17
Summary	19
Resources	19

Confluent Platform GitOps with ArgoCD

Complete Guide to Declarative Kafka Management on Kubernetes

Date: February 2, 2026

Author: Confluent Federal

Version: 2.0

Table of Contents

1. Executive Summary
2. Architecture Overview
3. Prerequisites
4. ArgoCD Installation & Configuration
5. Helm Charts Structure
6. Deploying Confluent Platform
7. Deploying Kafka Streams Applications
8. Deploying Flink Applications
9. Configuration Management
10. Branch-Based Deployments
11. Scaling Applications
12. Secrets Management
13. CI/CD Integration

-
14. Monitoring & Observability
 15. Operational Commands Reference
 16. Troubleshooting Guide
-

Executive Summary

This white paper documents a production-ready GitOps implementation for Confluent Platform on Kubernetes using ArgoCD. The solution provides:

- **Declarative Infrastructure** - All components defined as Helm charts
- **GitOps Deployment** - Git as the single source of truth
- **Modular Architecture** - Independent lifecycle for each application
- **Self-Healing** - ArgoCD automatically reverts configuration drift
- **Multi-Environment Support** - Branch and values-based configuration


Key Technologies

Technology	Purpose
Confluent Platform	Event streaming platform
CFK Operator	Kubernetes operator for Confluent
ArgoCD	GitOps continuous delivery
Helm	Package management and templating
Flink	Stream processing
Kafka Streams	Stateful stream processing

What This Guide Covers

- Installing and configuring ArgoCD
 - Deploying Confluent Platform components
 - Managing Kafka Streams applications (content routers)
 - Managing Flink applications
 - Environment-specific configurations
 - Monitoring, troubleshooting, and operations
-

Architecture Overview



- content-router - Apply manifests
- flink-application - Self-heal drift
- syslog-reconstruction
- datagen-connectors
- Schema Registry
- Kafka Connect
- ksqlDB
- Control Center
- Flink Jobs
- KStreams Apps

Repository Structure

```

cfk-argocd-demo/
  argocd/
    applications/                      # ArgoCD Application manifests
      confluent-platform-prod.yaml
      confluent-platform-dev.yaml
      content-router-prod.yaml
      content-router-syslog.yaml
      content-router-akamai.yaml
      flink-state-machine.yaml
      flink-kafka-streaming.yaml
      flink-hostname-enrichment.yaml
      syslog-reconstruction.yaml
      datagen-connectors-prod.yaml
    project.yaml                         # ArgoCD Project definition
  charts/
    confluent-platform/                 # Core platform chart
      Chart.yaml
      values.yaml                         # Default values
      values-dev.yaml                     # Dev overrides
      values-prod.yaml                   # Prod overrides
    templates/
      kafka.yaml
      kraftcontroller.yaml
      schemaregistry.yaml
      connect.yaml
      ksqldb.yaml
      controlcenter.yaml
      flink.yaml                          # Flink environment
    content-router/                      # KStreams routing app
      Chart.yaml
      values.yaml
      values-syslog.yaml
      values-akamai.yaml
      templates/
    flink-application/                  # Flink job chart
      Chart.yaml
      values.yaml
      values-state-machine.yaml
      values-kafka-streaming.yaml

```

```
values-hostname-enrichment.yaml  
templates/  
syslog-reconstruction/          # Specialized KStreams  
  Chart.yaml  
  values.yaml  
  templates/  
datagen-connectors/             # Connector definitions  
  Chart.yaml  
  values.yaml  
  templates/  
docs/  
  *.md                           # Documentation
```

Prerequisites

Required Tools

```
# Kubernetes CLI  
kubectl version --client  
  
# Helm  
helm version  
  
# Git  
git --version  
  
# GitHub CLI (optional, for repo management)  
gh --version
```

Kubernetes Cluster Requirements

- Kubernetes 1.25+
- Sufficient resources (minimum 3 nodes, 4 CPU / 16GB RAM each)
- Storage class configured (e.g., gp3, standard)
- CFK Operator installed

CFK Operator Installation

```
# Add Confluent Helm repository  
helm repo add confluentinc https://packages.confluent.io/helm  
helm repo update  
  
# Create namespace  
kubectl create namespace confluent-operator  
  
# Install CFK Operator  
helm upgrade --install confluent-operator \  
  confluentinc/confluent-for-kubernetes \  
  --namespace confluent-operator
```

```
--namespace confluent-operator \
--set namespaced=false \
--wait

# Verify CRDs
kubectl get crds | grep confluent
```

ArgoCD Installation & Configuration

Install ArgoCD

```
# Create namespace
kubectl create namespace argocd

# Install ArgoCD
kubectl apply -n argocd \
-f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml

# Wait for ArgoCD to be ready
kubectl wait --for=condition=ready pod \
-l app.kubernetes.io/name=argocd-server \
-n argocd \
--timeout=300s
```

Expose ArgoCD UI

```
# Expose via NodePort
kubectl patch svc argocd-server -n argocd \
-p '{"spec": {"type": "NodePort"}}

# Get access details
ARGOCD_PORT=$(kubectl get svc argocd-server -n argocd \
-o jsonpath='{.spec.ports[0].nodePort}')

NODE_IP=$(kubectl get nodes \
-o jsonpath='{.items[0].status.addresses[?(@.type=="ExternalIP")].address}')

ARGOCD_PASSWORD=$(kubectl -n argocd get secret argocd-initial-admin-secret \
-o jsonpath=".data.password" | base64 -d)

echo "URL: https://$NODE_IP:$ARGOCD_PORT"
echo "Username: admin"
echo "Password: $ARGOCD_PASSWORD"
```

Connect Repository

Via ArgoCD UI: 1. **Settings** → **Repositories** → + **Connect Repo** 2. Choose **VIA HTTPS** 3. Enter repository URL and GitHub credentials (PAT) 4. Click **Connect**

Via kubectl:

```
kubectl apply -f - <<EOF
apiVersion: v1
kind: Secret
metadata:
  name: repo-cfk-argocd-demo
  namespace: argocd
  labels:
    argocd.argoproj.io/secret-type: repository
stringData:
  type: git
  url: https://github.com/confluentfederal/cfk-argocd-demo.git
  username: <github-username>
  password: <github-pat>
EOF
```

Create ArgoCD Project

```
# argocd/project.yaml
apiVersion: argoproj.io/v1alpha1
kind: AppProject
metadata:
  name: confluent
  namespace: argocd
spec:
  description: Confluent Platform GitOps Project
  sourceRepos:
    - https://github.com/confluentfederal/cfk-argocd-demo.git
  destinations:
    - namespace: confluent
      server: https://kubernetes.default.svc
    - namespace: argocd
      server: https://kubernetes.default.svc
  clusterResourceWhitelist:
    - group: '*'
      kind: '*'
```

kubectl apply -f argocd/project.yaml

Helm Charts Structure

Chart Pattern

Each application follows a consistent Helm chart structure:

```
charts/<app-name>/
  Chart.yaml          # Chart metadata
  values.yaml         # Default configuration
```

```

values-<env>.yaml      # Environment-specific overrides
templates/
  _helpers.tpl        # Template helpers
  deployment.yaml     # Kubernetes Deployment
  configmap.yaml      # ConfigMap for app config
  ...                 # Additional resources

```

ArgoCD Application Pattern

Each ArgoCD Application references a Helm chart with specific values:

```

apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: <app-name>
  namespace: argocd
spec:
  project: confluent
  source:
    repoURL: https://github.com/confluentfederal/cfk-argocd-demo.git
    targetRevision: main
    path: charts/<chart-name>
    helm:
      releaseName: <release-name>
      valueFiles:
        - values-<env>.yaml
  destination:
    server: https://kubernetes.default.svc
    namespace: confluent
  syncPolicy:
    automated:
      prune: true
      selfHeal: true

```

Deploying Confluent Platform

Chart: confluent-platform

The core platform chart deploys: - KRaft Controllers - Kafka Brokers - Schema Registry - Kafka Connect - ksqlDB - Control Center - Flink Environment (CMF + Operator)

Deploy Production Environment

```
kubectl apply -f argocd/applications/confluent-platform-prod.yaml
```

ArgoCD Application:

```

apiVersion: argoproj.io/v1alpha1
kind: Application

```

```

metadata:
  name: confluent-platform-prod
  namespace: argocd
spec:
  project: confluent
  source:
    repoURL: https://github.com/confluentfederal/cfk-argocd-demo.git
    targetRevision: helm-migration
    path: charts/confluent-platform
    helm:
      releaseName: confluent
      valueFiles:
        - values.yaml
        - values-prod.yaml
  destination:
    server: https://kubernetes.default.svc
    namespace: confluent
  syncPolicy:
    automated:
      prune: true
      selfHeal: true

```

Verify Deployment

```

# Check ArgoCD status
kubectl get application confluent-platform-prod -n argocd

# Check Confluent components
kubectl get kafka,connect,schemaregistry,ksqldb,controlcenter -n confluent

# Check pods
kubectl get pods -n confluent

```

Deploying Kafka Streams Applications

Chart: content-router

A reusable chart for Kafka Streams content routing applications.

Pattern: Multiple Instances from One Chart

Each content router instance uses the same chart with different values:

Instance	Values File	Purpose
content-router-prod	values.yaml	Default routing
content-router-syslog	values-syslog.yaml	Syslog routing
content-router-akamai	values-akamai.yaml	Akamai CDN routing

Deploy a Content Router

```
kubectl apply -f argocd/applications/content-router-syslog.yaml
```

ArgoCD Application:

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: content-router-syslog
  namespace: argocd
spec:
  project: confluent
  source:
    repoURL: https://github.com/confluentfederal/cfk-argocd-demo.git
    targetRevision: helm-migration
    path: charts/content-router
    helm:
      releaseName: content-router-syslog
      valueFiles:
        - values-syslog.yaml
  destination:
    server: https://kubernetes.default.svc
    namespace: confluent
  syncPolicy:
    automated:
      prune: true
      selfHeal: true
```

Adding a New Content Router Instance

1. Create values file: `charts/content-router/values-<name>.yaml`
 2. Create ArgoCD application: `argocd/applications/content-router-<name>.yaml`
 3. Commit and push to Git
 4. Apply the ArgoCD application
-

Deploying Flink Applications

Chart: flink-application

A reusable chart for deploying Flink jobs via FlinkApplication CRD.

Prerequisites

The Confluent Platform chart creates: - CMFRestClass (Confluent Manager for Flink) - FlinkEnvironment (`flink-env`)

Deploy a Flink Application

```
kubectl apply -f argocd/applications/flink-kafka-streaming.yaml
```

Values file example:

```
# charts/flink-application/values-kafka-streaming.yaml
namespace: confluent
flinkEnvironment: flink-env

image:
  repository: cstevenson954/java-flink
  tag: "0.0.3"

flinkVersion: v1_20
serviceAccount: flink

flinkConfiguration:
  taskmanager.numberOfTaskSlots: "2"

jobManager:
  resource:
    memory: "1024m"
    cpu: "0.5"

taskManager:
  resource:
    memory: "1024m"
    cpu: "0.5"

job:
  jarURI: local:///opt/flink/usrlib/java-flink-0.0.1.jar
  entryClass: com.example.javaflink.StreamingJob
  state: running
  parallelism: 1
  upgradeMode: stateless
  args:
    - "--bootstrap.servers=kafka.confluent.svc.cluster.local:9071"
    - "--input.topic=flink-input"
    - "--output.topic=flink-output"
```

Verify Flink Deployment

```
# Check FlinkApplication
kubectl get flinkapplication -n confluent

# Check FlinkDeployment (created by CMF)
kubectl get flinkdeployments -n confluent

# Check pods
kubectl get pods -n confluent | grep flink
```

Configuration Management

Values File Hierarchy

1. `values.yaml` - Base defaults
2. `values-<env>.yaml` - Environment overrides (merged on top)

Example: Environment-Specific Resources

`values.yaml` (defaults):

```
kafka:  
  replicas: 3  
  resources:  
    requests:  
      cpu: "1"  
      memory: "4Gi"
```

`values-dev.yaml` (dev overrides):

```
kafka:  
  replicas: 1  
  resources:  
    requests:  
      cpu: "500m"  
      memory: "2Gi"
```

`values-prod.yaml` (prod overrides):

```
kafka:  
  replicas: 3  
  resources:  
    requests:  
      cpu: "2"  
      memory: "8Gi"
```

Applying Configuration Changes

1. Edit the values file in Git
2. Commit and push
3. ArgoCD auto-syncs (or force refresh):

```
kubectl annotate application <app-name> -n argocd \  
argocd.argoproj.io/refresh=hard --overwrite
```

Branch-Based Deployments

Strategy: targetRevision

Use Git branches to manage different environments or feature deployments.

Branch	Purpose
main	Production-ready code
develop	Integration testing
feature/*	Feature development
helm-migration	Helm chart development

Switching Branches

Via ArgoCD UI: 1. Click application → **App Details** → **Edit** 2. Change **Target Revision** to desired branch 3. Save

Via kubectl:

```
kubectl patch application confluent-platform-prod -n argocd \
--type merge -p '{"spec":{"source":{"targetRevision":"develop"}}}'
```

Multi-Environment Setup

Deploy same chart to different environments using different ArgoCD Applications:

```
# Dev environment
spec:
  source:
    targetRevision: develop
    helm:
      valueFiles:
        - values-dev.yaml
```

```
# Prod environment
spec:
  source:
    targetRevision: main
    helm:
      valueFiles:
        - values-prod.yaml
```

Scaling Applications

Scaling Kafka Streams Apps

Update `replicas` in values file:

```
# charts/content-router/values-syslog.yaml
replicas: 4 # Scale from 2 to 4
```

Commit, push, and ArgoCD applies the change.

Scaling Flink Jobs

Update parallelism and resources:

```
# charts/flink-application/values-kafka-streaming.yaml
job:
  parallelism: 4

taskManager:
  resource:
    memory: "2048m"
    cpu: "1"
```

Scaling Kafka Brokers

Update broker count in values:

```
# charts/confluent-platform/values-prod.yaml
kafka:
  replicas: 5
```

Note: Scaling down Kafka requires partition reassignment.

Secrets Management

Option 1: Kubernetes Secrets (Basic)

Create secrets manually, reference in values:

```
kubectl create secret generic kafka-credentials \
  -n confluent \
  --from-literal=username=admin \
  --from-literal=password=secret

# In values.yaml
kafka:
  authentication:
    secretRef: kafka-credentials
```

Option 2: External Secrets Operator

Integrate with AWS Secrets Manager, HashiCorp Vault, etc.:

```
apiVersion: external-secrets.io/v1beta1
kind: ExternalSecret
metadata:
  name: kafka-credentials
  namespace: confluent
spec:
  secretStoreRef:
    name: aws-secrets-manager
```

```

kind: ClusterSecretStore
target:
  name: kafka-credentials
data:
  - secretKey: username
    remoteRef:
      key: confluent/kafka
      property: username
  - secretKey: password
    remoteRef:
      key: confluent/kafka
      property: password

```

Option 3: Sealed Secrets

Encrypt secrets in Git:

```

# Install kubeseal CLI
# Encrypt secret
kubeseal --format yaml < secret.yaml > sealed-secret.yaml

```

Store `sealed-secret.yaml` in Git; SealedSecrets controller decrypts at runtime.

CI/CD Integration

GitHub Actions Example

```

# .github/workflows/deploy.yml
name: Deploy to ArgoCD

on:
  push:
    branches: [main]
    paths:
      - 'charts/**'

jobs:
  sync:
    runs-on: ubuntu-latest
    steps:
      - name: Trigger ArgoCD Sync
        run: |
          curl -X POST \
            -H "Authorization: Bearer ${{ secrets.ARGOCD_TOKEN }}" \
            "https://${{ secrets.ARGOCD_SERVER }}/api/v1/applications/confluent-platform-prod/"


```

ArgoCD Webhook Integration

Configure GitHub webhook to notify ArgoCD of changes:

1. ArgoCD Settings → Repositories → Configure webhook
2. Add webhook URL to GitHub repository settings
3. ArgoCD immediately syncs on push

GitOps Workflow



Monitoring & Observability

ArgoCD Status Monitoring

```

# List all applications with status
kubectl get applications -n argocd

# Watch application status
kubectl get application <name> -n argocd -w

# Get sync details
kubectl get application <name> -n argocd \
-o jsonpath='{.status.sync.status} / {.status.health.status}'
  
```

Control Center Access

```

# Get Control Center URL
NODE_IP=$(kubectl get nodes -o jsonpath='{.items[0].status.addresses[?(@.type=="ExternalIP")].address}')
CC_PORT=$(kubectl get svc controlcenter -n confluent -o jsonpath='{.spec.ports[0].nodePort}')
echo "http://$NODE_IP:$CC_PORT"
  
```

Application Logs

```

# Kafka Streams app logs
kubectl logs -l app=content-router-syslog -n confluent --tail=100

# Flink job manager logs
kubectl logs -l app=kafka-streaming-job,component=jobmanager -n confluent

# Flink task manager logs
kubectl logs -l app=kafka-streaming-job,component=taskmanager -n confluent
  
```

```
# Connect worker logs
kubectl logs -l app=connect -n confluent --tail=100
```

Flink Dashboard

```
# Port-forward to Flink UI
kubectl port-forward svc/kafka-streaming-job-rest 8081:8081 -n confluent
# Access at http://localhost:8081
```

Metrics & Alerting

For production, integrate with: - **Prometheus** - Metrics collection - **Grafana** - Dashboards - **AlertManager** - Alerting

CFK exposes metrics endpoints on all components.

Operational Commands Reference

ArgoCD Commands

Command	Description
kubectl get applications -n argocd	List all applications
kubectl get application <name> -n argocd -o yaml	View application details
kubectl annotate application <name> -n argocd argocd.argoproj.io/refresh=hard	Force sync
kubectl delete application <name> -n argocd	Delete application

Confluent Platform Commands

Command	Description
kubectl get kafka,connect,schemaregistry -n confluent	List CP components
kubectl get connector -n confluent	List connectors
kubectl get flinkapplication -n confluent	List Flink apps
kubectl get flinkdeployments -n confluent	List Flink deployments

Kafka Topic Management

Command	Description
kubectl exec kafka-0 -n confluent -- kafka-topics --list --bootstrap-server localhost:9092	List topics
kubectl exec kafka-0 -n confluent -- kafka-topics --create --topic <name> --partitions 3 --replication-factor 3 --bootstrap-server localhost:9092	Create topic
kubectl exec kafka-0 -n confluent -- kafka-console-consumer --topic <name> --bootstrap-server localhost:9092 --from-beginning	Consume messages

Troubleshooting Guide

ArgoCD Sync Issues

Application stuck in “Syncing”:

```
# Force refresh
kubectl annotate application <name> -n argocd \
    argocd.argoproj.io/refresh=hard --overwrite
```

```
# Clear stuck operation
kubectl patch application <name> -n argocd \
    --type merge -p '{"operation": null}'
```

Application “OutOfSync”:

```
# Check diff
kubectl get application <name> -n argocd -o jsonpath='{.status.sync.revision}'

# Force sync
kubectl patch application <name> -n argocd \
    --type merge -p '{"operation": {"initiatedBy": {"username": "admin"}, "sync": {}}}'
```

Flink Application Issues

FlinkApplication stuck in CREATED:

```

# Check CMF logs
kubectl logs deployment/confluent-manager-for-apache-flink -n confluent --tail=50

# Restart CMF
kubectl rollout restart deployment/confluent-manager-for-apache-flink -n confluent

```

FlinkDeployment FAILED:

```

# Get error details
kubectl describe flinkdeployment <name> -n confluent | grep -A10 "Error:"

```

```

# Check job manager logs
kubectl logs -l app=<name>,component=jobmanager -n confluent

```

Common Flink Errors:

Error	Cause	Fix
NoClassDefFoundError	Missing dependency in JAR	Fix Maven shade plugin order
Insufficient cpu	Cluster resource limits	Reduce resource requests
Memory configuration failed	Memory too low	Increase to minimum 1024m

Kafka Streams Issues

Pod CrashLoopBackOff:

```

# Check logs
kubectl logs <pod-name> -n confluent --previous

# Common causes:
# - Missing Kafka topics
# - Invalid bootstrap servers
# - Configuration errors

```

Create missing topics:

```

kubectl exec kafka-0 -n confluent -- kafka-topics \
--create --topic <name> \
--partitions 3 --replication-factor 3 \
--bootstrap-server localhost:9092

```

Resource Issues

Pods Pending (Insufficient CPU/Memory):

```

# Check node resources
kubectl describe nodes | grep -A 10 "Allocated resources"

# Reduce application resources in values.yaml

```

```
resources:  
  requests:  
    cpu: "100m"  
    memory: "512Mi"
```

Summary

This GitOps implementation provides:

	Benefit	Description
Single Source of Truth		All configuration in Git
Audit Trail		Git history tracks all changes
Rollback		Revert to any previous state
Self-Healing		ArgoCD reverts manual drift
Multi-Environment		Values files and branches
Modular		Independent app lifecycles

Current Deployments

Application Type	Count	Chart
Confluent Platform	1	confluent-platform
Content Routers	3	content-router
Flink Applications	4	flink-application
Syslog Reconstruction	1	syslog-reconstruction

Resources

- **Repository:** <https://github.com/confluentfederal/cfk-argocd-demo>
- **CFK Documentation:** <https://docs.confluent.io/operator/current/overview.html>
- **ArgoCD Documentation:** <https://argo-cd.readthedocs.io/>
- **Flink on Confluent:** <https://docs.confluent.io/platform/current/flink/overview.html>

Document Version: 2.0 / Last Updated: February 2, 2026